

m1-l2-p1

January 28, 2024

1 M1-L2 Problem 1 (10 points)

In this problem, we are given a function $L(w_1, w_2)$ with a known functional form. You will perform gradient descent to find a global minimum. The goal is to find what initial guesses and learning rates (step sizes) lead the algorithm to find the global minimum.

The function $L(w_1, w_2)$ is defined as:

$$L(w_1, w_2) = \cos(4w_1 + w_2/4 - 1) + w_2^2 + 2w_1^2$$

A Python function for $L(w_1, w_2)$ is given.

1.1 Gradients

First, we must define a gradient of L . That is $\nabla L = \left[\frac{\partial L}{\partial w_1}, \frac{\partial L}{\partial w_2} \right]$. First, compute these derivatives by hand. Then, in the cell below, complete the functions for the derivatives of L with respect to w_1 and w_2 .

```
[4]: import numpy as np
import matplotlib.pyplot as plt

def L(w1, w2):
    return np.cos(4*w1 + w2/4 - 1) + w2*w2 + 2*w1*w1

def dLdw1(w1, w2):
    # YOUR CODE GOES HERE
    return -4*np.sin(4*w1 + w2/4 - 1) + 4*w1

def dLdw2(w1, w2):
    # YOUR CODE GOES HERE
    return -1/4*np.sin(4*w1 + w2/4 - 1) + 2*w2
```

1.2 Gradient Descent

The function `plot_gd` performs gradient descent by calling your derivative functions. Take a look at how this works. Then, run the interactive gradient descent cell that follows and answer the questions below.

```
[5]: def plot_gd(w1, w2, log_stepsize, log_steps):
    stepsize = 10**log_stepsize
```

```

steps = int(10**log_steps)

# Gradient Descent
w1s = np.zeros(steps+1)
w2s = np.zeros(steps+1)

for i in range(steps):
    w1s[i], w2s[i] = w1, w2
    w1 = w1 - stepsize * dLdw1(w1s[i],w2s[i])
    w2 = w2 - stepsize * dLdw2(w1s[i],w2s[i])
w1s[steps], w2s[steps] = w1, w2

# Plotting
vals = np.linspace(-1,1,50)
x, y = np.meshgrid(vals,vals)
z = L(x,y)

plt.figure(figsize=(7,5.8),dpi=120)
plt.contour(x,y,z,colors="black", levels=np.linspace(-.5,3,6))
plt.pcolormesh(x,y,z,shading="nearest",cmap="Blues")
plt.colorbar()

plt.plot(w1s,w2s,"g-",marker=".
↪",markerfacecolor="black",markeredgcolor="None")
plt.scatter(w1s[0],w2s[0],zorder=100, color="blue",marker="o",label=f"$w_0$")
↪= [{w1s[0]:.1f}, {w2s[0]:.1f}])
plt.scatter(w1,w2,zorder=100,color="red",marker="x",label=f"$w^*$ = [{w1:.
↪2f}, {w2:.2f}]")
plt.legend(loc="upper left")

plt.axis("equal")
plt.box(False)
plt.xlabel("$w_1$")
plt.ylabel("$w_2$")
plt.xlim(-1,1)
plt.ylim(-1,1)
plt.title(f"Step size = {stepsize:.0e}; {steps} steps")
plt.show()

```

```

[6]: %matplotlib inline
from ipywidgets import interact, interactive, fixed, interact_manual, Layout,
↪FloatSlider, Dropdown

slider1 = FloatSlider(
    value=0,
    min=-1,
    max=1,

```

```

        step=.1,
        description='w1 guess',
        disabled=False,
        continuous_update=True,
        orientation='horizontal',
        readout=False,
        layout = Layout(width='550px')
    )

slider2 = FloatSlider(
    value=0,
    min=-1,
    max=1,
    step=.1,
    description='w2 guess',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

slider3 = FloatSlider(
    value=-1.5,
    min=-3,
    max=0,
    step=.5,
    description='step size',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

slider4 = FloatSlider(
    value=2,
    min=0,
    max=3,
    step=.25,
    description='steps',
    disabled=False,
    continuous_update=True,
    orientation='horizontal',
    readout=False,
    layout = Layout(width='550px')
)

```

```

interactive_plot = interactive(
    plot_gd,
    w1 = slider1,
    w2 = slider2,
    log_stepsize = slider3,
    log_steps = slider4,
)
output = interactive_plot.children[-1]
output.layout.height = '620px'

interactive_plot

```

```

[6]: interactive(children=(FloatSlider(value=0.0, description='w1 guess',
    layout=Layout(width='550px'), max=1.0, mi...

```

1.3 Questions

Play around with the sliders above to get an intuition for which initial conditions/learning rates lead us to find the global minimum at $[-0.42, -0.05]$. Then answer the following questions:

1. Set w_0 to $[0.2, 0.8]$ and step size to $1e-01$. After 100 steps of gradient descent, what w^* do we reach? With the given parameters, w^* reached $[-0.42, -0.05]$.
2. Keep parameters from the previous question, but change the initial guess to $[0.3, 0.8]$. Now what is the optimum we find? With the given parameters, the optimum is at $[0.80, 0.10]$.
3. Set w_0 to $[-1.0, -1.0]$ and number of iterations to 1000 and step size to $1e-03$. What w^* do we reach, and why is it not exactly the global minimum? With the given parameters, w^* reached $[-0.42, -0.18]$, which is most likely a local minimum. This is a result of step size being too small and the model can not converge to the global minimum with just 1000 steps.
4. In general, what happens if we set learning rate too large? A large learning rate will lead the model to overshoot in the direction of the gradient, which can cause oscillations around the minimum and not converge. It can also lead to model instability and poor learning performance.