

m2-l2-p3

February 3, 2024

1 M2-L2 Problem 3 (5 points)

When flow is directed across a pin fin heat sink, increasing fluid velocity can improve the heat transfer, making the heat sink more effective.

You have been given a dataset containing measurements for such a scenario, which contains the following: - Input: Reynolds Number of air flowing past the heat sink - Output: Heat transfer coefficient of the heat sink, in $\text{W}/(\text{m}^2 \text{ K})$

Your job is to train a model on this data to predict the heat transfer coefficient, given Reynolds number as input. You will use a high-order polynomial

Start by loading the data in the following cell:

```
[5]: import numpy as np
import matplotlib.pyplot as plt

def plot_data_with_regression(x_data, y_data, x_reg, y_reg, title=""):
    plt.figure()

    plt.scatter(x_data.flatten(), y_data.flatten(), label="Data", c="black")
    plt.plot(x_reg.flatten(), y_reg.flatten(), label="Fit")

    plt.legend(loc="upper left")
    plt.xlabel(r"$Re / 1000$")
    plt.ylabel(r"$h, \text{W}/\text{m}^2 \text{ K}$")
    plt.xlim(0,6)
    plt.ylim(50,200)
    plt.title(title)
    plt.show()


deg = 5
x = np.array([1.010, 2.000, 2.990, 4.100, 5.020])
y = np.array([75.1, 104.0, 100.6, 138.8, 150.8])
X = np.vander(x, deg)

xreg = np.linspace(0,6)
Xreg = np.vander(xreg, deg)
```

1.1 Least Squares Regression

As we have done for previous problems, we can do least squares regression by computing the pseudo-inverse of the design matrix. Notice how the model performs beyond the training data.

```
[6]: w = np.linalg.inv(X.T @ X) @ X.T @ y.reshape(-1,1)
yreg = Xreg @ w
plot_data_with_regression(x, y, xreg, yreg, "5th order polynomial regression")
```



1.2 L2 Regularization

Notice that the plot above reveals that our fifth-order model is overfitting to the data. Let's try applying L2 regularization to fix this. In the lecture, the closed-form solution to least squares with L2 regularization was:

$$w = (X'X + \lambda I_m)^{-1} X'y$$

where I_m is the identity matrix, but with zero in the bias row/column instead of 1; λ is regularization strength; X' is the design matrix and y column vector output.

Complete the function below to compute this w for a given λ :


```
[7]: def get_regularized_w(L):
    I_m = np.eye(deg)
    I_m[-1,-1] = 0


    # YOUR CODE GOES HERE
    # return regularized w
    regularized_w = np.linalg.inv(X.T @ X + L*I_m) @ X.T @ y.reshape(-1,1)
    return regularized_w
```


1.3 Testing different lambda values


With the above function written, we can compute w for some different values of λ and decide which is qualitatively best.

```
[8]: for L in [0,.001,0.1,1000]:
    w = get_regularized_w(L)
    yreg = Xreg @ w
    plot_data_with_regression(x, y, xreg, yreg,
                             f"Regression with Regularization,  $\lambda$ ={L}")
```









1.4 Model Selection

Which value of lambda appears to yield the “best” model?

When lambda is 0.1, the model appears to yield the best.